

# Foresight Lens

Resource exhaustion prediction system for Tier 1 cloud services.

- Lead time  $\geq 15$  min
- Inference ~\$0 cost
- Model ewma-stl-v1




Fig. 1 — EWMA control chart issuing early alerts ~106 minutes before SLO breach.



# When the monitor turns **red**, it's already too late

Traditional monitoring only triggers at 99% — leaving SREs with zero response time.

# 8 Core Requirements

- **Proactive prediction:** Must alert before actual system crash (Lead time  $\geq$  15 mins).
- **High Precision:** FP rate  $\leq$  12%, detecting  $\geq$  80% of drifts.
- **Cost Control:** Total monthly AWS bill strictly under \$200.
- **Multi-tenant Security:** Applied to 3 tier-1 services with isolated baselines and zero cross-tenant bleed.
- **No Auto-remediation:** AI strictly acts as an advisor (PURELY Predictive) without system intervention.
- **Per-service Baseline:** Train baselines per service (1 manual run) + ADRs defining retrain triggers.
- **Actionable Recommendation:** Must include 5 components (action verb, target, from  $\rightarrow$  to, confidence, evidence).
- **Audit Log:** Log every prediction request, encrypted at rest.

$\leq$  \$200

AWS BUDGET/MONTH

Advise

NO AUTO-REMEDIATION

3 svc

PAYMENT · FRAUD · LEDGER

# Out of Scope (Client Confirmed)

Strict boundaries established from W11 to ensure focus and operational safety.

## Functional Exclusions

- **Auto-remediation:** Predict & recommend only.
- **Cross-service RCA:** Per-service drift only.
- **Cost forecasting:** Transferred to Task Force 2.
- **Business Metrics:** Infra metrics only (no fraud rate, tx count).

## Operational Exclusions

- **Auto-retrain pipeline:** Design ADR is sufficient.
- **Multi-region deployment:** Single region, DR design-only.
- **Production traffic mirror:** Local load test with k6/Locust.
- **Scale targets:** 99.5% SLO (not 99.99%) & exactly 3 Tier-1 services.



# Traditional monitoring alerts **too late**

Failure symptoms persist for over 90 minutes, yet systems only alert at the absolute brink.

## 4 Factors Driving the Shift to Predictive Alerts

### 1. Silent Exhaustion

Incidents don't explode instantly; they silently decay (Memory leak, queue backlog) over 90+ minutes.

### 2. Alert Fatigue

A low Static threshold causes endless False Positives; a high one only alerts post-crash.

### 3. 24/7 Blindspot

Systems generate too many metrics. Nobody has the time to stare at 12 Grafana dashboards 24/7.

### 4. Dynamic Seasonality

Fintech traffic has stark day/night cycles. Relying on a fixed baseline is entirely meaningless.

HARSH REALITY

# Illustration: Silent Exhaustion Incident

The system struggles for 3 continuous hours, yet the Static threshold remains completely blind.




Fig. 2 — A memory leak creeping up silently for 180 minutes. The 99% mark (Static threshold) flashes red only upon crashing. Lead time = 0.



# Early predictions, clear advice — **no unauthorized actions**

Strictly acting as an Advisor. The CDO platform retains ultimate decision-making authority.



MEASURABLE OUTCOMES (1/4)

# 106 Minutes Lead Time


Memory leak anomaly alerted 106 minutes before capacity exhaustion crash.  
Exceeds the  $\geq 15\text{m}$  requirement.

```
// evidence_lead_time.json

{
  "method": "Simulate gradual CPU drift → detect via EWMA...",
  ...
  "detected_at_minute": 734,
  "breach_at_minute": 840,
  "lead_time_minutes": 106,
  "requirement": "≥15 minutes",
  "pass": true
}
```

# 3 Isolated Tenants

Per-service baselines strictly maintain tenant isolation with statistical proof of pattern diversity.



```
// evidence_pattern_diversity.json

{
  "claim": "3 services exhibit distinctly different CPU patterns",
  "tests": {
    "anova": { "f_statistic": 3234.51, "p_value": 0.0, "significant": true },
    "peak_hour_analysis": {
      "payment-gw": 10,
      "ledger": 1,
      "fraud-detector": 13
    }
  },
  "conclusion": "Static thresholds would trigger massive False Positives."
}
```

EVIDENCE REFERENCE  
evidence\_pattern\_diversity.json

MEASURABLE OUTCOMES (3/4)

# \$0 Token Cost

In-house statistical model costs zero LLM tokens. Complete AWS footprint runs at ~\$53/month.

```
// evidence_cost.json

{
  "services": {
    "ECS Fargate": { "cost": "$36.00/month" },
    "Application Load Balancer": { "cost": "$16.00/month" },
    "API_Gateway_private": { "cost": "$0" },
    "S3_audit_and_baselines": { "cost": "$1.00/month" }
  },
  "total_estimated_monthly": "$53.00",
  "budget_cap": "$200/month",
  "conclusion": "Using ECS Fargate is cost-effective (~$53/mo)... Inference token cost is $0. "
}
```

# 5 Action Components

Alerts are bundled with specific action verbs, targets, transitions, confidence scores, and evidence.

```
// evidence_recommendation.json

{
  "requirement_met": "action verb + target + from→to + confidence + evidence link",
  "test_cases": [
    {
      "metric": "cpu_usage_percent",
      "recommendation": "SCALE_UP (Action) RDS Database (Target) from db.r5.large to db.r5.xlarge (From→To).  
Confidence: 0.85 (runtime-computed). Evidence: https://obs.internal/tnt-1?metric=cpu_usage_percent"
    },
    ...
  ]
}
```

05

METHODOLOGY & RATIONALE

# Why we built it this way

Behind every design decision is a rigorous defense backed by hard evidence.

# Synthetic Baseline Generation

How we generated statistically rigorous training data and holdout sets for the STL model.

## 7-Day Telemetry Simulation

To train our models without access to production data, we simulated 7 days of 1-minute granularity telemetry per service.

- **Base Seasonal Trend:** `daily_shape()` applies a "business-hours double-hump load curve" (peaking at 9–11 AM and 2–4 PM).
- **Noise Injection:** Gaussian white noise `RNG.normal(0, sigma)` added to simulate real-world traffic fluctuations.
- **STL Extraction:** We then ran STL (period=1440) on the clean data (Days 0–5) to extract the seasonal profile and residual standard deviation.

## Holdout Set & Honest Evaluation

We isolated **Day 6** as our Holdout evaluation set and injected 4 specific anomalies to measure detection speed:

- **Gradual Drift:** CPU ramps from 40% to 94% over 90 mins.
- **Sudden Spike:** Latency jumps +600ms for 20 mins.
- **Slow Leak:** Memory creeps from 50% to 96% over 180 mins.
- **Noisy FP Trap:** Massive queue depth variance but *no real mean shift* (Detector must stay silent).

Evidence: `scripts/train_baseline.py`

# Grounding tests in Historical Outages

We didn't invent random test cases. We mapped our Multi-Root-Cause Injection framework directly to the Client's exact pain points from the last 3 months.

CLIENT BRIEF (W11)

"RDS CPU bò lên 100% trong 90 phút..."

Test Vector 1

CPU Exhaustion

CLIENT BRIEF (W11)

"...trước khi connection pool exhaust."

Test Vector 2

Connection Pool Leak

CLIENT BRIEF (W11)

"Queue worker backlog âm thâm 6x..."

Test Vector 3

Queue Backlog

CLIENT BRIEF (W11)

"Memory leak dẫn đến OOM..."

Test Vector 4

Memory Leak

EVIDENCE REFERENCE

TF4\_FORESIGHT\_LEARNER.md (Client Brief W11)

# Visualizing the Lagging Trap

Across all simulated historical outages, **API Latency** and **Queue Depth** consistently break bounds long before CPU or Memory.




Fig 1 — Detection timeline. Latency detects at minute 722, Memory only flags at 781.



# Metric Priority Assessment

Ranking the metrics based on detection speed across all fault vectors.




Fig 2 — Overall Metric Priority Score.

# Mathematical Proof of Priority

We aggregated detection ranks across all 4 root causes. The result completely ruled out CPU and Memory as viable early-warning signals.

```
// evidence_metric_priority.json

{
  "method": "Multi-root-cause injection: 4 different failure modes",
  "circular_reasoning_fix": "Previous method injected only CPU drift. Now inject 4 independent root causes.",
  "root_causes": [
    { "root_cause": "CPU exhaustion" },
    { "root_cause": "Memory leak" },
    { "root_cause": "Queue backlog" },
    { "root_cause": "Connection pool leak" }
  ],
  "average_rank": {
    "api_latency_ms": 1.0,
    "queue_depth": 2.0,
    "cpu_usage_percent": 3.25,
    "memory_usage_percent": 3.75
  },
  "final_rank": [
    "api_latency_ms", "queue_depth",
    "cpu_usage_percent", "memory_usage_percent"
  ],
  "conclusion": "Across 4 root causes, api_latency_ms is #1 in 4/4 root causes."
}
```

EVIDENCE REFERENCE  
evidence\_metric\_priority.json

# De-seasonalise: STL Decomposition

If we use static thresholds, daytime traffic peaks will trigger false alarms.




## Stripping the Day/Night Cycle

Infrastructure metrics naturally rise during business hours (e.g., 09:00–11:00) and fall at night. A simple threshold would constantly fire false alerts during peak load.

**STL (Seasonal and Trend decomposition using Loess)** mathematically extracts this daily pattern offline. When running in production, we subtract the expected pattern from the live signal, leaving only the "raw residual truth" to be scanned for anomalies.


EVIDENCE REFERENCE

evidence\_seasonality.json



# Detect: The EWMA Radar

A memory leak isn't a sudden spike; it's a slow, creeping death.



## Catching Sustained Drift

Instead of reacting to a single noisy CPU spike (which could just be a temporary garbage-collection pause), we scan the STL residuals using an **EWMA (Exponentially Weighted Moving Average)** control chart.

EWMA acts as a low-pass filter. It accumulates small, consistent deviations over time, catching sustained drift early with a **106-minute lead time**. Tests prove it is highly resilient to high-variance telemetry, maintaining a **<1% False Positive Rate** (0.004) under severe Gaussian noise (std=20).

LEAD TIME PROOF

evidence\_lead\_time.json

NOISE ROBUSTNESS

evidence\_noisy\_baseline.json

# Why Statistical won over Machine Learning

We evaluated Isolation Forest (ML) against STL+EWMA on our holdout data. The results forced our hand.


## The False Positive Problem

Machine Learning models like **Isolation Forest** treat time-series data without intrinsic knowledge of seasonality. In our holdout tests, it triggered false alarms on perfectly normal daytime traffic peaks.

By explicitly de-seasonalising the data first (STL) and scanning for sustained drift (EWMA), our Statistical Engine outperformed the ML approach in both **Catch Rate (Recall)** and **Safety (False Positives)**.

EVIDENCE REFERENCE

evidence\_algorithm\_comparison.json



ISOLATION FOREST

21.4% FP

Failed ≤12% gate

STL + EWMA

7.1% FP

Safely passed

# An Advisor, Not an Autonomous Actor

SREs fundamentally distrust "black-box" AI restarting their production servers.

We deliberately chose to abandon "Auto-Remediation". Foresight Lens acts strictly as a **Predictive Advisor**. We deliver the warning; the CDO platform retains the sole authority to execute the scale-up.

To enforce this, we implemented strict **Confidence Gating**. If the engine's Brier-calibrated confidence score drops below 0.7, the recommendation is forcefully downgraded to **INVESTIGATE**.

**Evidence:** `evidence_confidence_threshold.json`

- **Operating Point:** EWMA  $\alpha=0.3$ ,  $K=4.0$
- **Brier Score:** **0.049** (*Score < 0.25 indicates excellent probability calibration*)
- **Result:** We only issue a **SCALE\_UP** command when the mathematical probability of exhaustion is overwhelmingly high, preserving SRE trust.




ENGINEERING & VALIDATION

# Architected for **production trust**

Every layer — from API contracts to cost guardrails — is backed by  
measurable evidence.

# Unidirectional Data Flow & Stateless Processing


Microservices strictly decoupled via HTTP API. Latency < 5ms via NumPy in-memory vectorization.





# Automated Data & Deployment Pipelines

Ensuring reproducible training, stateless inference, and serverless containerized deployments.



# Strict API Boundaries & Stable Interfaces

Protecting system integrity by separating telemetry collection from decision reasoning.

## 01 · Telemetry Contract

```
// 120-min rolling window
{
  "tenant_id": "payment-core",
  "resolution": "1-minute",
  "gauges": ["cpu", "mem"]
}
```

**Imputation strictly enforced.** Missing data yields immediate HTTP 400 Bad Request.

## 02 · AI API (Additive v1.1)

```
// POST /v1/predict
{
  "sla": "p99 < 500ms",
  "responses": {
    "200": "Recommendation",
    "422": "Schema Violation",
    "503": "Fail Open"
  }
}
```

**Unit Normalization:** Handled unit discrepancies (e.g., ms vs us) at the boundary via Pydantic mapping.


## 03 · Deployment Contract

```
// aws_ecs.yml
{
  "compute": "Fargate",
  "scaling": "min 2 / max 4",
  "circuit_breaker": {
    "max_cost": "$200/mo"
  }
}
```

**Stateless Resilience:** Deployed on multi-AZ ECS Fargate for high availability and automatic recovery.

# Statistical Rigor & Holdout Evaluation


Measured on 790 windows across 3 tier-1 services with 4 injected drift scenarios + FP traps.



PERFORMANCE VALIDATION

# Load Test: 400 RPS with Zero Errors

We stress-tested the engine at 4x the SLA target. Result: p99 latency under 5ms, 100% success rate.



```
// evidence_load_test.json (excerpt)

{
  "phase1_sla": {
    "target_rps": 100, "requests_sent": 3000,
    "success_rate": 1.0,
    "latency_ms": { "p50": 2.78, "p95": 3.53, "p99": 4.01 }
  },
  "phase2_capacity_probe": {
    "target_rps": 400, "requests_sent": 6000,
    "success_rate": 1.0,
    "latency_ms": { "p50": 1.7, "p95": 1.98, "p99": 2.71 }
  },
  "overall_pass": true
}
```

EVIDENCE REFERENCE  
evidence\_load\_test.json

# Flat-Rate Serverless Economics

Rejecting LLM per-token pricing in favor of in-house statistical compute.

\$53 / month

Total TCO — 26.5% of the \$200 limit.

\$0

INFERENCE TOKENS

< \$1

MARGINAL COST / TENANT

## Cost Breakdown

- **ECS Fargate (~\$36):** Flat billing by task-second, NOT request volume. Min 2 / Max 4 tasks auto-scale.
- **Internal ALB (~\$16):** Base hourly + LCU.
- **Private API Gateway (~\$0):** Within AWS free tier for demo volumes.
- **S3 & KMS (~\$1):** Baseline JSON storage and audit log encryption.

**Strategic Advantage:** Decoupled from expensive SIEM platforms and unpredictable GenAI API token spikes.

# 28–Minute Service Onboarding

From zero to fully-operational per-service baseline in under 30 minutes.



```
// evidence_onboarding.json

{
  "steps": {
    "register": "~5 minutes (manual input + config gen)",
    "load_7d_data": "~10 minutes (read 30K data points from S3/TSDB)",
    "train_baseline": "~10 minutes (STL decomposition + EWMA baseline)",
    "validate": "~3 minutes (run test predictions + check results)"
  },
  "total_estimated": "28-30 minutes per service"
}
```

EVIDENCE REFERENCE

evidence\_onboarding.json

# Security-by-Design & Defense in Depth

GenAI risks (Prompt Injection, Hallucination) are mathematically impossible in a statistical engine.

### 01 · Deterministic Output

Same JSON in, same output out.  
Reproducible mathematical execution with zero hallucination.

### 02 · Confidence Gating

Hardcoded safety logic: `if confidence < 0.7` → downgrade to INVESTIGATE.

### 03 · Traceable Audit

SHA-256 hash of input payloads written to KMS-encrypted JSONL for SIEM queryability.

### 04 · Fail-Open Degrade

Engine outage triggers CDO fallback to rules. Stateless design allows instant ECS restart (**RTO < 60s**).

### 05 · Context Isolation

Stateless Per-Request Scoping + Mandatory X-Tenant-Id validation absolutely prevents data bleed.

### 06 · DoS Protection

Pydantic enforces max 10,000 datapoints. API Gateway applies 600 req/min/tenant limit.

TRADE-OFFS

# Architecture Decision Records (ADRs)

Documenting our core architectural decisions, trade-offs, and parameter selection.

ADR-001: Statistical vs LLM

**Chosen:** STL + EWMA.  
**Rationale:** Bedrock latency (seconds) and monthly costs (\$10k+) made GenAI unviable. Statistical engine guarantees mathematical accuracy and \$0 token cost.

ADR-002: Confidence Threshold

**Chosen:** 0.7 Gate.  
**Rationale:** Balances recall with Alert Fatigue. Weak signals (<0.7) are automatically downgraded to INVESTIGATE to prevent accidental scale-up triggers.

ADR-003: Flash Sales Silence

**Chosen:** Manual Silence.  
**Rationale:** Unpredictable spikes during planned campaigns (e.g. Black Friday) will trigger alerts. SREs can manually silence/retrain baselines for these events.

ADR-004: STL vs Isolation Forest

**Chosen:** STL + EWMA.  
**Rationale:** ML model (Isolation Forest) yielded 21.4% FP rate on holdout data. De-seasonalising first using STL reduced the FP rate to 7.1%.

ADR-005: Baseline Refresh

**Chosen:** Weekly Refresh.  
**Rationale:** Baseline trained manually using offline script. Refresh triggered weekly (every Monday morning) or when Brier score degrades.

ADR-006: EWMA Tuning

**Chosen:**  $\alpha=0.3$ ,  $K=4.0$ .  
**Rationale:** 16-point sweep optimized detection.  $K=3.0$  triggered 17.5% FP rate (failed gate);  $K=4.0$  kept FP at 7.1% while maintaining 97.1% recall.



## ACHIEVEMENTS

# Surpassed all stringent Client requirements, **exceeding expectations**

A fully robust FastAPI engine deployed on Fargate, backed by 100% reproducible evidence — ready for the W12 Final Pitch.

~\$53

26.5% BUDGET

3

EXECUTED  
CONTRACTS

9/9

TESTS COMPLETELY  
PASSED

106  
min

LEAD TIME

## Backup Slides (For Q&A)

### EWMA Parameter Tuning

| A    | K   | FP%  | RECALL |
|------|-----|------|--------|
| 0.3  | 3.0 | 17.5 | .977   |
| 0.3  | 4.0 | 7.1  | .971   |
| 0.3  | 4.5 | 6.5  | .966   |
| 0.15 | 4.0 | 11.8 | .977   |

### Telemetry · 7 Signals

- cpu\_usage\_percent
- memory\_usage\_percent
- api\_latency\_ms
- queue\_depth
- active\_connections
- db\_connection\_pool\_pct
- cache\_hit\_rate\_pct

4 utilizing baseline · 3 falling back to z-score

### API Error Codes

| CODE | DESCRIPTION                               |
|------|-------------------------------------------|
| 400  | Bad request — Invalid payload format      |
| 401  | Unauthorized — Missing API key            |
| 422  | Validation error — Invalid schema         |
| 429  | Rate limited — Exceeded tenant quota      |
| 503  | Service unavailable — Fail-open triggered |